

DevOps Handbook

Part I

- CH 1 - Agile, continuous delivery, and the three ways
- CH 2 - The first way: the principles of flow
- CH 3 - The second way: the principles of feedback
- CH 4 - The third way: the principles of continual learning and experimentation

Part II

- CH 5 - Selecting which value stream to start with pg. 51
 - Greenfield vs. brownfield services pg. 54
 - Consider both **systems of record** (ERP, HR, Finance) and **systems of engagement** (employee or customer facing) pg. 56
 - Start with the most sympathetic and innovative groups
 - Technology adoption curve Fig 9. Pg. 58
 - Expanding DevOps across our organization pg. 58
 - PH 1. Find innovators and early adopters
 - PH 2. Build critical mass and silent majority
 - PH 3. Identify the holdouts
 - “Little fish learn to be big fish in little ponds.” - Peter Drucker
- CH 6 - Understanding the work in our value stream, making it visible, and expanding it across the organization pg. 61
 - Identify the teams supporting our value stream pg. 63
 - Create a value stream map to see the work pg. 63
 - Creating a dedicating transformation team pg. 66
 - Agree on a shared goal pg. 68
 - Keep our improvement planning horizons short pg.68
 - Reserve 20% of cycles for **non-functional requirements** and reducing technical debt pg. 69
 - Positive user-invisible value Figure 11 pg. 70
 - Increase the visibility of work pg. 73
 - Use tools to reinforce desired behavior pg. 73
- CH 7 - How to design our organization and architecture with Conway’s Law in mind pg. 77
 - ORMs pg. 79
 - Organizational archetypes pg. 80
 - Functional, Matrix, Market
 - Problems often caused by overly functional orientation (“optimized for cost”) pg. 81
 - PSC, IT, HR, shared services
 - Enable market-oriented teams (“optimized for speed”) pg. 82

- Making functional orientation work pg. 83
 - Testing, operations, and security as everyone's job, every day pg. 84
 - Enable every team member to be a generalist pg. 85
 - Specialists vs. generalists vs. "E-shaped" Table 2 pg. 86
 - QE transition last yr
 - Dr. Carol Dweck fixed vs. growth mindset pg.87
 - Fund not projects, but services and products pg. 87
 - Design team boundaries in accordance with Conway's Law pg. 88
 - Create loosely-coupled architectures to enable developer productivity and safety pg. 89
 - Keep team sizes small (the "two-pizza team" rule) pg. 90
 - CH 8 - How to get great outcomes by integrating operations into the daily work of development pg. 95
 - Create shared services to increase developer productivity pg. 97
 - Embed ops engineers into our service teams pg. 99
 - Assign an ops engineer to each service team pg. 100
 - Integrate ops into dev rituals pg. 101
 - Invite ops to our dev standups pg. 102
 - Invite ops to our dev retrospectives pg. 102
 - Make relevant ops work visible on shared kanban boards pg. 104
-

Part III

- CH 9 - Create the Foundations of Our Deployment Pipeline
 - Enable on-demand creation of dev, test, and production environments pg. 113
 - Use automation for any or all:
 - Copying a virtualized environment (eg. VMware image, etc.)
 - Building an automated environment creation process that starts from “bare metal” (eg, PXE install from a baseline image)
 - Using “infrastructure as code” config management tools (eg Puppet, Chef, Ansible, Salt, CFEngine, etc.)
 - Using automated operating system configuration tools (eg Solaris Jumpstart, Red Hat Kickstart, Debian preseed)
 - Assembling an environment from a set of virtual images or containers (eg. Vagrant, Docker)
 - AWS, Azure, other public cloud
 - Create our single repository of truth for the entire system pg. 115
 - All application code and dependencies
 - Any script used to create db schemas, application reference data
 - All env create tools
 - Any file used to create containers
 - All supporting automated tests and any manual test scripts
 - Any script that supports code packaging
 - Cloud config files
 - *****Use of version control by operations is the highest predictor of both IT performance and organizational performance***** pg 117
 - **Make infrastructure easier to rebuild than to repair** pg. 118
 - Pets vs. Cattle pg 118 - Bill Baker (Microsoft) quote
 - Make changes to one config, then automatically deploy everywhere via Puppet or Ansible, etc.
 - *Immutable infrastructure*
 - Modify our definition of “done” to include running in production-like environments
- CH 10 - Enable fast and reliable automated testing
 - *Imposter syndrome* pg. 124
 - *Fix-it days - improvement blitzes* pg 125
 - *Containers* pg. 128
 - *Test pyramid* pg. 133
 - *Test Driven Development* pg. 134
- CH 11 - Enable and practice continuous integration
 - *Version control branches* pg. 143
 - *Trunk-based development* pg. 145
 - *Optimize for indiv productivity*
 - *Optimize for team productivity*

- *BLUF don't use feature branching pg. 147 AKA small code batch sizes*
 - CH 12 - Automate and enable low-risk releases
 - *Fully document the steps in the deployment process with an eye towards automating as many as possible. Pg. 155*
 - *Mean Time To Repair (MTTR) pg. 158*
 - *Decouple deployments from releases pg. 164*
 - *Environment-based release patterns pg. 165*
 - *Blue-green deployment pattern pg. 166*
 - *Database changes pg. 167*
 - *Application-based release patterns pg. 165*
 - *Dark launches*
 - CH 13 - Architect for low-risk releases
 - *Architect for low-risk releases*
 - *Strangler application pattern pg. 180*
 - *Martin Fowler*
-

Part IV

- CH 14 - Create telemetry to enable seeing and solving problems
 - *Create telemetry to enable seeing and solving problems*
 - *Create within our applications and environments*
 - *Ganglia, Graphite*
 - *Mean Time To Repair (MTTR) in minutes not days*
 - *Graph on pg. 197*
 - *Art of Monitoring by James Turnball*
 - *Example tools footnote pg. 199*
 - *Get logs, transform them into metrics using the event router*
 - *“Monitoring is so important that our monitoring systems need to be more available and scalable than the systems being monitored.” - Adrian Cockcroft*
 - *logging of critical features*
 - *INFO – wrong login for FPY*
 - *WARN – DB not writing or taking too long to connect*
 - *WARN – Internet connectivity*
 - *inability to create institutional knowledge*
 - *80% of all outages are caused by change and 80% of MTTR is spent trying to determine what changed*
 - *questions on pg. 204*
 - *Enable creation of Production metrics as part of daily work*
 - *StatsD from Etsy*
 - *Create self-service access to telemetry and **information radiators***
 - ***IDEA** – create a 320F new product launch dashboard*
 - *When metrics aren’t actionable, they are likely vanity metrics that provide little useful information*
 - *infrastructure metrics*
 - *cost of delayed features*
 - *instead of measuring downtime, measure real business consequences of downtime. How much revenue should we have attained but didn’t.*
- CH 15 - Analyze telemetry to better anticipate problems and achieve goals
 - *Analyze telemetry to better anticipate problems and achieve goals*
 - *outlier detection*
 - *Normally distributed data*
 - *compute mean and std for normal*
 - *what to alert on*
 - *analyze severe outages in last 30 days*
 - *create a list of telemetry that could have enabled earlier and faster detection and diagnosis*
 - *anomaly detection*
 - *chi-squared distribution*
 - *smoothing - moving averages (rolling avgs)*

- *Fast Fourier Transform test*
 - *Kolmogorov Smirnov and other "non-parametric"*
 - *R*
 - *Oculus*
 - *Opsweekly*
 - *Skyline*
 - *Application logs*
 - *INFO - wrong FPY tool login*
 - *WARN - Db writes not working or taking too long*
 - *WARN - internet connectivity or uptime*
 - CH 16 - Enable feedback so development and operations can safely deploy code
 - *Pg 228 Progression*
 - *Optimize for MTTR instead of MTBF*
 - *Fix forward or roll back*
 - *Contextual inquiry pg 232*
 - *UX observation*
 - *Service handback mechanism*
 - *Fig. 38 Pg 237*
 - *Launch Readiness Review & Handoff Readiness Review*
 - CH 17 - Integrate hypothesis-driven development and A/B testing into our daily work
 - *Balsamiq*
 - *A/B testing*
 - CH 18 - Create review and coordination processes to increase quality of our current work
 - *Peer review*
 - *Pull request*
 - *Gitflow pg 250*
 - *"The Knight Capital" failure pg 251*
 - *Counterfactual*
 - *Code review*
 - *"Ask a programmer to review ten lines of code, he'll find ten issues. Ask him to do five hundred lines, and he'll say it looks good." pg 256*
 - *Pair programming pg 256*
 - *Pair pattern: driver/ observer*
 - *Pair pattern: one write tests/ other implement*
 - *Extreme programming pg 259*
-

Part V

- CH 19 - Enable and inject learning into daily work
 - *Dr. Steven Spears quote pg. 271*
 - *NETFLIX “Cloud Native” architecture pg. 272*
 - *Chaos Monkey*
 - *Establish a just, learning culture*
 - *Dr. Sidney Dekker*
 - *Bad Apple Theory*
 - *“Human error is not the cause of troubles; instead, human error is a consequence of the design of the tools that we gave them”*
 - *Instead of “naming, blaming, and shaming or goal should always be to maximize opportunities for organizational learning.*
 - *Blameless postmortems pg. 274-275*
 - *Controlled introduction of failures into production pg. 274*
 - *John Allspaw quote pg. 274*
 - *System as imagined vs system that actually exists*
 - *Publish post mortems as widely as possible*
 - *Search: “Chef” “postmortem”*
 - CH 20 - Convert local discoveries into global improvements
 - *Chatbots and chat rooms pg. 287*
 - *Automate standardized processes in software for re-use pg. 289*
 - *Create a single, shared source code repository for our entire organization pg. 290*
 - *Spread knowledge by using automated tests as documentation and communities of practice pg. 293*
 - *Design for operations through codified nonfunctional requirements pg. 293*
 - *Build reusable operations user stories into development pg. 294*
 - *Ops checklist pg. 295*
 - *Ensure technology choices help achieve organizational goals pg. 295*
 - CH 21 - Reserve time to create organizational learning and improvement
 - *Institutionalize rituals to pay down technical debt pg. 300*
 - *Enable everyone to teach and learn pg. 303*
 - *Create internal consulting and coaches to spread practices pg. 306*
-

Part VI

- CH 22 - Information security as everyone's job, every day
 - *James Wickett creator of Gauntlet security tool*
 - *Integrate security into development iteration demonstrations*
 - *Possible metrics: development velocity, failed customer interactions*
 - *Integrate security into defect tracking and post-mortems*
 - *Integrate preventive security controls into shared source code repositories and shared services*
 - *Integrate security into our deployment pipeline*
 - *Ensure security of the application*
 - CH 23 - Protecting the deployment pipeline
 - *Integrate security and compliance into change approval processes*
 - *Re-categorize the majority of our lower risk changes as standard changes*
 - *What to do when changes are categorized as normal changes*
 - *Reduce reliance on separation of duty*
 - *Ensure documentation and proof for auditors and compliance officers*
-